

```
lemp001@u20-05:~$ sudo apt-get install openjdk-8-jdk
[sudo] password for lemp001:
Sorry, try again.
[sudo] password for lemp001:
Sorry, try again.
[sudo] password for lemp001:
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following packages were automatically installed and are no longer required:
  chromium-codecs-ffmpeg-extra gstreamer1.0-vaapi libfwupdplugin1
  libgstreamer-plugins-bad1.0-0 libva-wayland2
Use 'sudo apt autoremove' to remove them.
The following additional packages will be installed:
  ca-certificates-java fonts-dejavu-extra java-common libatk-wrapper-java
  libatk-wrapper-java-jni libice-dev libpthread-stubs0-dev libsm-dev
  libx11-dev libxau-dev libxcb1-dev libxdmcp-dev libxt-dev
  openjdk-8-jdk-headless openjdk-8-jre openjdk-8-jre-headless
  x11proto-core-dev x11proto-dev xorg-sgml-doctools xtrans-dev
Suggested packages:
  default-jre libice-doc libsm-doc libx11-doc libxcb-doc libxt-doc
  openjdk-8-demo openjdk-8-source visualvm fonts-ipafont-gothic
  fonts-ipafont-mincho fonts-wqy-microhei fonts-wqy-zenhei
The following NEW packages will be installed:
```

```
lemp001@u20-05:~$ java -version
openjdk version "1.8.0_362"
OpenJDK Runtime Environment (build 1.8.0_362-8u362-ga-0ubuntu1~20.04.1-b09)
OpenJDK 64-Bit Server VM (build 25.362-b09, mixed mode)
```

```
lemp001@u20-05:~$ sudo nano /etc/environment
lemp001@u20-05:~$ sudo adduser hadoop
Adding user `hadoop' ...
Adding new group `hadoop' (1001) ...
Adding new user `hadoop' (1001) with group `hadoop' ...
Creating home directory `/home/hadoop' ...
Copying files from `/etc/skel' ...
New password:
Retype new password:
passwd: password updated successfully
Changing the user information for hadoop
Enter the new value, or press ENTER for the default
    Full Name []:
    Room Number []:
    Work Phone []:
    Home Phone []:
    Other []:
Is the information correct? [Y/n] y
```

172.26.35.25 — Подключение к удаленному рабочему столу

anp 21 13:~

lemp001@u20-05: ~

GNU nano 4.8 /etc/environment Modified

```
PATH="/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin:/usr/games:/"
JAVA_HOME="/usr/lib/jvm/java-8-openjdk-amd64"
JRE_HOME="/usr/lib/jvm/java-8-openjdk-amd64/jre"
```

File Name to Write: /etc/environment

^G Get Help	M-D DOS Format	M-A Append	M-B Backup File
^C Cancel	M-M Mac Format	M-P Prepend	^T To Files

lemp001@u20-05: ~

GNU nano 4.8 /etc/environment

```
PATH="/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin:/usr/games:/"
JAVA_HOME="/usr/lib/jvm/java-8-openjdk-amd64"
JRE_HOME="/usr/lib/jvm/java-8-openjdk-amd64/jre"
```

[Wrote 3 lines]

^G Get Help	^O Write Out	^W Where Is	^K Cut Text	^J Justify	^C Cur Pos
^X Exit	^R Read File	^\ Replace	^U Paste Text	^T To Spell	^_ Go To Line

```
lemp001@u20-05:~$ sudo su hadoop
hadoop@u20-05:/home/lemp001$ cd
hadoop@u20-05:~$ ssh-keygen -t rsa -N "" -f /home/hadoop/.ssh/id_rsa
Generating public/private rsa key pair.
Created directory '/home/hadoop/.ssh'.
Your identification has been saved in /home/hadoop/.ssh/id_rsa
Your public key has been saved in /home/hadoop/.ssh/id_rsa.pub
The key fingerprint is:
SHA256:L3DUzsZBt30coxNP7iq/pfAvPl5CGGw1IdBmMV4h7Bw hadoop@u20-05
The key's randomart image is:
+---[RSA 3072]-----+
|      .+*.Bo      |
|      =EB + .     |
|      .B.+ *      |
|      . =0+ 0 .    |
|      . S * + o    |
|      o o . o      |
|      . o . +      |
|      ..oo*        |
|      =0=.         |
+----[SHA256]-----+
hadoop@u20-05:~$ cat /home/hadoop/.ssh/id_rsa.pub >> /home/hadoop/.ssh/authorized_keys
hadoop@u20-05:~$ chmod 0600 /home/hadoop/.ssh/authorized_keys
hadoop@u20-05:~$
```

Welcome to Ubuntu 20.04.4 LTS (GNU/Linux 5.13.0-51-generic x86_64)

```
* Documentation:  https://help.ubuntu.com
* Management:    https://landscape.canonical.com
* Support:       https://ubuntu.com/advantage
```

87 updates can be applied immediately.
1 of these updates is a standard security update.
To see these additional updates run: apt list --upgradable

Your Hardware Enablement Stack (HWE) is supported until April 2025.
*** System restart required ***

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

```
hadoop@u20-05:~$ exit
logout
Connection to localhost closed.
```

```
hadoop@u20-05:~$ wget https://archive.apache.org/dist/hadoop/common/hadoop-3.1.2/hadoop-3.1.2.tar.gz
--2023-04-21 13:51:40-- https://archive.apache.org/dist/hadoop/common/hadoop-3.1.2/hadoop-3.1.2.tar.gz
Resolving archive.apache.org (archive.apache.org)... 138.201.131.134, 2a01:4f8:172:2ec5::2
Connecting to archive.apache.org (archive.apache.org)|138.201.131.134|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 332433589 (317M) [application/x-gzip]
Saving to: 'hadoop-3.1.2.tar.gz'

hadoop-3.1.2.tar.gz      100%[=====] 317,03M   710KB/s   in 7m 50s

2023-04-21 13:59:31 (690 KB/s) - 'hadoop-3.1.2.tar.gz' saved [332433589/332433589]
```

```
hadoop@u20-05: ~
GNU nano 4.8 .bashrc
# ~/.bashrc: executed by bash(1) for non-login shells.
# see /usr/share/doc/bash/examples/startup-files (in the package bash-doc)
# for examples

# If not running interactively, don't do anything
case $- in
  *i*) ;;
  *) return;;
esac

# don't put duplicate lines or lines starting with space in the history.
# See bash(1) for more options
HISTCONTROL=ignoreboth

# append to the history file, don't overwrite it
shopt -s histappend

# for setting history length see HISTSIZE and HISTFILESIZE in bash(1)
HISTSIZE=1000
HISTFILESIZE=2000

# check the window size after each command and, if necessary,
[ Read 117 lines ]
^G Get Help  ^O Write Out  ^W Where Is  ^K Cut Text   ^J Justify    ^C Cur Pos    M-U Undo
^X Exit      ^R Read File  ^_ Replace   ^U Paste Text ^T To Spell   ^_ Go To Line  M-E Redo

hadoop@u20-05:~$ mv hadoop-3.1.2 hadoop
hadoop@u20-05:~$ nano .bashrc
hadoop@u20-05:~$ source .bashrc
hadoop@u20-05:~$
```

```
hadoop@u20-05: ~
GNU nano 4.8 /home/hadoop/hadoop/etc/hadoop/hadoop-env.sh Modified
# and therefore may override any similar flags set in HADOOP_OPTS
#
# export HDFS_DFSROUTER_OPTS=""
###
###
# Advanced Users Only!
###
#
# When building Hadoop, one can add the class paths to the commands
# via this special env var:
# export HADOOP_ENABLE_BUILD_PATHS="true"
#
# To prevent accidents, shell commands be (superficially) locked
# to only allow certain users to execute certain subcommands.
# It uses the format of (command)_(subcommand)_USER.
#
# For example, to limit who can execute the namenode command,
# export HDFS_NAMENODE_USER=hdfs
export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64
^G Get Help  ^O Write Out  ^W Where Is  ^K Cut Text   ^J Justify    ^C Cur Pos    M-U Undo
^X Exit      ^R Read File  ^_ Replace   ^U Paste Text ^T To Spell   ^_ Go To Line  M-E Redo
```

```
hadoop@u20-05: ~
GNU nano 4.8 /home/hadoop/hadoop/etc/hadoop/core-site.xml
You may obtain a copy of the License at

http://www.apache.org/licenses/LICENSE-2.0

Unless required by applicable law or agreed to in writing, software
distributed under the License is distributed on an "AS IS" BASIS,
WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
See the License for the specific language governing permissions and
limitations under the License. See accompanying LICENSE file.
-->

<!-- Put site-specific property overrides in this file. -->

<configuration>
<property>
<name>fs.default.name</name>
<value>hdfs://localhost:9000</value>
</property>
</configuration>

[ Wrote 25 lines ]
^G Get Help ^O Write Out ^W Where Is ^K Cut Text ^J Justify ^C Cur Pos M-U Undo
^X Exit ^R Read File ^\ Replace ^U Paste Text ^T To Spell ^_ Go To Line M-E Redo
```

```
hadoop@u20-05: ~
GNU nano 4.8 /home/hadoop/hadoop/etc/hadoop/hdfs-site.xml Modified
-->

<!-- Put site-specific property overrides in this file. -->

<configuration>
<property>

<name>dfs.replication</name>

<value>1</value>

</property>

<property>

<name>dfs.name.dir</name>

<value>file:///home/hadoop/hadoopdata/hdfs/namenode</value>
</property>
<property>
<name>dfs.data.dir</name>
<value>file:///home/hadoop/hadoopdata/hdfs/datanode</value>

^G Get Help ^O Write Out ^W Where Is ^K Cut Text ^J Justify ^C Cur Pos M-U Undo
^X Exit ^R Read File ^\ Replace ^U Paste Text ^T To Spell ^_ Go To Line M-E Redo
```

```
hadoop@u20-05: ~
<configuration>
<property>

<name>mapreduce.framework.name</name>

<value>yarn</value>

</property>

<property>

<name>yarn.app.mapreduce.am.env</name>

<value>HADOOP_MAPRED_HOME=${HADOOP_HOME}</value>

</property>

<property>

<name>mapreduce.map.env</name>

<value>HADOOP_MAPRED_HOME=${HADOOP_HOME}</value>

</property>

```

```
hadoop@u20-05: ~
GNU nano 4.8 /home/hadoop/hadoop/etc/hadoop/mapred-site.xml

<value>HADOOP_MAPRED_HOME=${HADOOP_HOME}</value>
</property>
<property>
<name>mapreduce.map.env</name>
<value>HADOOP_MAPRED_HOME=${HADOOP_HOME}</value>
</property>
<property>
<name>mapreduce.reduce.env</name>
<value>HADOOP_MAPRED_HOME=${HADOOP_HOME}</value>
</property>
</configuration>

```

^G Get Help	^O Write Out	^W Where Is	^K Cut Text	^J Justify	^C Cur Pos	M-U Undo
^X Exit	^R Read File	^_ Replace	^U Paste Text	^T To Spell	^ Go To Line	M-E Redo

```
hadoop@u20-05: ~
GNU nano 4.8 /home/hadoop/hadoop/etc/hadoop/yarn-site.xml Modified
WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
See the License for the specific language governing permissions and
limitations under the License. See accompanying LICENSE file.
-->
<configuration>

<!-- Site specific YARN configuration properties -->
<property>

<name>yarn.nodemanager.aux-services</name>

<value>mapreduce_shuffle</value>

</property>

<property>

<name>yarn.nodemanager.resource.memory-mb</name>

<value>16384</value>

</property>
File Name to Write: /home/hadoop/hadoop/etc/hadoop/yarn-site.xml
^G Get Help      M-D DOS Format   M-A Append      M-B Backup File
^C Cancel        M-M Mac Format   M-P Prepend     ^T To Files
```

```
hadoop@u20-05:~$ hadoop fs -mkdir /user
2023-04-21 14:35:21,671 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... u
sing builtin-java classes where applicable
hadoop@u20-05:~$ hadoop fs -mkdir /user/hadoop
2023-04-21 14:35:36,377 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... u
sing builtin-java classes where applicable
hadoop@u20-05:~$ hadoop fs -put /var/log/dpkg.log /user/hadoop/dpkg.log
2023-04-21 14:37:15,148 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... u
sing builtin-java classes where applicable
hadoop@u20-05:~$
```

All Applications x Browsing HDFS x +

localhost:9870/explorer.html#/user/hadoop

Hadoop Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities

Browse Directory

/user/hadoop

Go!

Show 25 entries

Search:

<input type="checkbox"/>	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
<input type="checkbox"/>	-rw-r--r--	hadoop	supergroup	314.76 KB	Apr 21 14:37	1	128 MB	dpkg.log	

Showing 1 to 1 of 1 entries

Previous

1

Next

Hadoop, 2018.

```
hadoop@u20-05:~$ ' hadoop fs -cat /user/hadoop/test_output/part-r-00000
2023-04-21 14:44:23,238 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform..
asses where applicable
0.10.2-0ubuntu0.20.04.1 4
0.11.0-0ubuntu0.20.04.1 6
0.136ubuntu6.7 12
0.17-2 48
0.18.2-2 4
0.18.2-2ubuntu0.1 6
0.19.5-1ubuntu1 4
0.19.5-1ubuntu1.1 6
0.24-1ubuntu3 68
0.36-6ubuntu1 4
0.37.1-1 14
0.38.4-0ubuntu1 4
0.38.4-0ubuntu2.1 6
0.4-1 7
0.72 7
0.86.1-0ubuntu1 16
0.86.1-0ubuntu1.1 24
0.9.12-1 7
06:14:20 7
06:14:21 8
06:14:25 5
06:14:26 5
06:14:27 13
06:14:28 2
06:14:28 2
```

Show 20 ▾ entries

ID ▾	User ↕	Name ↕	Application Type ↕	Queue ↕	Application Priority ↕	StartTime ↕	FinishTime ↕	State ↕	FinalStatus ↕	Running Containers ↕	Allocated Containers
application_1682076545743_0001	hadoop	word count	MAPREDUCE	default	0	Fri Apr 21 14:40:56 +0300 2023	Fri Apr 21 14:41:18 +0300 2023	FINISHED	SUCCEEDED	N/A	N/A

Showing 1 to 1 of 1 entries